

Quiz 5

Deadline	Friday, 24 April 2020 at 11:59PM
Latest Submission	Saturday, 18 April 2020 at 7:55PM
Maximum Mark	4

Question 1 (1 mark)

Consider a table of machine parts (*Parts(id,name,colour,size,cost)*), about which we have the following information:

```
db=# select count(distinct size) from Parts;
count
-----
      5

db=# select min(size) from Parts;
min
-----
tiny

db=# select max(size) from Parts;
max
-----
huge

db=# select count(*) from Parts;
count
-----
6507
```

Approximately how many tuples would you expect to find in the result of the following query?

```
select * from Parts where size = 'huge';
```

Assume a uniform distribution of part sizes.

(a) <input type="radio"/>	650
(b) <input type="radio"/>	3253
(c) <input type="radio"/>	6506
(d) <input type="radio"/>	None of the other options is correct.

(e) <input checked="" type="radio"/>	1301
--------------------------------------	------

Question 2 (1 mark)

Consider the three relations $R(id,a,b)$, $S(rid,tid,c)$, $T(id,d,e,f)$ where the id attributes are primary keys, and $S.rid$ is a foreign key referencing $R.id$ and $S.tid$ is a foreign key referencing $T.id$.

Consider also the following join on these three relations:

```
select a,b,c,d,e,f from R, S, T where R.id = S.rid and T.id = S.tid
```

Which of the following does **not** represent a possible join order for the above SQL statement?

(a) <input type="radio"/>	$(R \bowtie S) \bowtie T$
(b) <input type="radio"/>	$R \bowtie (S \bowtie T)$
(c) <input type="radio"/>	$(S \bowtie T) \bowtie R$
(d) <input checked="" type="radio"/>	$S \bowtie (T \bowtie R)$
(e) <input type="radio"/>	They are all valid join orders for the query.

Question 3 (1 mark)

Consider the following two relations/tables:

$R(a,b,c)$ where $r_R = 100000$, $c_R = 100$, $b_R = 1000$

$S(c,d,e)$ where $r_S = 50000$, $c_S = 500$, $b_S = 100$

If we do a natural join on these two tables, using a block nested loop join with 35 buffers, how many pages do we *read* in completing the join?

(a) <input type="radio"/>	1100
(b) <input type="radio"/>	100000
(c) <input type="radio"/>	None of the other options is correct
(d) <input type="radio"/>	3100
(e) <input checked="" type="radio"/>	4100


Question 4 (1 mark)

Consider two tables $R(id,x,y,z)$ and $S(id,a,b,rid)$ where rid is a foreign key referring to $R.id$. Consider also the following join on these tables:

```
select x,a from R, S where R.id = S.rid;
```

This join is implemented using the hybrid hash join algorithm with 17 memory buffers, 1 buffer used for input (#0), 8 buffers used to hold one partition in memory (#1-#8), 7 output buffers used to transfer tuples in the other partitions to disk (#9-#16), and 1 "spare" output buffer (#16) used only in the first phase (see below).

The diagram below shows the first phase of the hash join, where R is scanned and partitioned into 8 hash buckets. Partition 0 is stored in memory; partitions 1-7 are written to disk.



After the first phase has partitioned R , the join algorithm then partitions S using the same hash function. Any tuples that hash to partition 0 are matched against R 's partition 0 tuples held in memory buffers 1-8, and any resulting matches are written to disk via the spare output buffer. Other S tuples are written to disk-based partitions. Partitions 1-7 for R and S are then processed as for a standard hash join, using a second hash function.

Assume that:

- we have a well-behaved (uniform) hash functions for both the first and second phases
- R contains 3000 tuples in 60 pages; each partition of R requires exactly 8 pages
- S contains 1600 tuples in 40 pages; each partition of S requires exactly 5 pages
- each tuple in S refers to a different tuple in R (i.e. $S.rid$ is unique)
- the result contains 1600 tuples which are written into 30 pages (i.e. they count as disk I/O)

Based on the above, compute the number of disk I/Os needed to execute this join.

(a) <input type="radio"/>	308
(b) <input type="radio"/>	272
(c) <input checked="" type="radio"/>	312
(d) <input type="radio"/>	338
(e) <input type="radio"/>	None of the other options is correct

✓ Submit